

A Prime Curvature Hamiltonian on the Logarithmic Axis with 0.657% Agreement to the Riemann Spectrum

Jack Pickett- Bank House, Cornwall — 13th December 2025

Abstract

A one-dimensional, self-adjoint Hamiltonian on the logarithmic prime axis,

$$\hat{H} = -\frac{d^2}{dt^2} + V(t), \quad t = \log p,$$

where the potential $V(t)$ is derived directly from a curvature field built from local composite densities around the primes. On a log-uniform grid generated from the first 10^6 primes, the lowest 80 eigenvalues of the discretised operator are compared to the first 80 non-trivial zeros γ_n of the Riemann zeta function.

Using a three-parameter affine-log model

$$\gamma_n \approx a \lambda_n + c \log n + b$$

fitted on the first 20 levels and evaluated on the first 80, we obtain for the 10^6 prime run:

- mean relative error **0.657%**,
- maximum relative error **2.892%**

The same construction with 2×10^5 primes yields a mean error of 2.463% and a maximum error 5.773%. Nearest-neighbour spacing statistics of the unfolded spectrum show strong level repulsion and are clearly incompatible with a Poisson process, though more rigid than the Gaussian Unitary Ensemble (GUE). We present this as a concrete semiclassical Hilbert–Pólya candidate: a real, log-scale Hamiltonian whose low spectrum tracks the Riemann zeros at the one-percent level.

1. Overview

The Hilbert–Pólya idea is that the non-trivial zeros of the Riemann zeta function

$$\zeta\left(\frac{1}{2}+iy_n\right) = 0$$

should arise as eigenvalues of a self-adjoint operator on an appropriate Hilbert space. If such an operator \hat{H} exists with

$$\text{spec}(\hat{H}) = \{\gamma_n\}$$

then the Riemann Hypothesis follows from self-adjointness.

A **specific, explicitly defined** operator built from the primes reproduces the first part of the Riemann spectrum with high quantitative accuracy, using no complex analysis and no fitted “ad hoc” potential. The only inputs are:

- the list of primes p
- local composite densities in fixed windows around the primes,
- a small number of global scaling parameters.

Method:

1. Build a **curvature field** $\kappa(p)$ from composite densities in windows around each prime.
2. Interpolate this onto a log-uniform grid $t = \log p$ to obtain κ
3. Construct a Schrödinger-type operator:

$$\hat{H} = -\frac{d^2}{dt^2} + V(t), \quad V(t) = \beta \kappa(t) + \varepsilon \log t,$$

and discretise it as a tridiagonal matrix on the log grid.

4. Compute its lowest eigenvalues λ_n and compare to the first Riemann zeros γ_n via a simple three-parameter regression.

The result is a **numerical operator model** tied directly to the primes, with spectral output closely tracking the Riemann zeros.

2. Curvature field from the primes

2.1 Prime sample and composite density

Fix a number of primes N_p . In this work we use:

- a calibration run with $N_p = 200\,000$
- a main run with $N_p = 1\,000\,000$

We generate the first N_p primes by a standard sieve up to a safe upper bound. Let $\{p_k\}_{k=1}^{N_p}$ denote this list. Around each prime p we consider the integer window:

$$[p - R, p + R], \quad R = 20.$$

Using a sieve up to $\max p + R$, we construct a boolean array marking primes and composites.

The **local composite density** $\rho(p)$ is the fraction of composite numbers in this window:

$$\rho(p) = \frac{\#\{n \in [p - R, p + R] : n \text{ composite}\}}{2R + 1}$$

2.2 Curvature definition

From the density we define a local scale factor:

$$\sigma(p) = \log(1 + \rho(p) \log p)$$

and a curvature field:

$$\kappa(p) = C \sigma(p)^3 \sqrt{\rho(p)}$$

with a fixed constant $C = 0.150$. The cubic dependence on σ and the square root in ρ are empirical but fixed once chosen; there is no tuning per prime.

We obtain a dataset:

$$\{(p_k, \kappa_k)\}_{k=1}^{N_\kappa}$$

where some boundary primes near 2 may be dropped if the window would exit the sieved domain.

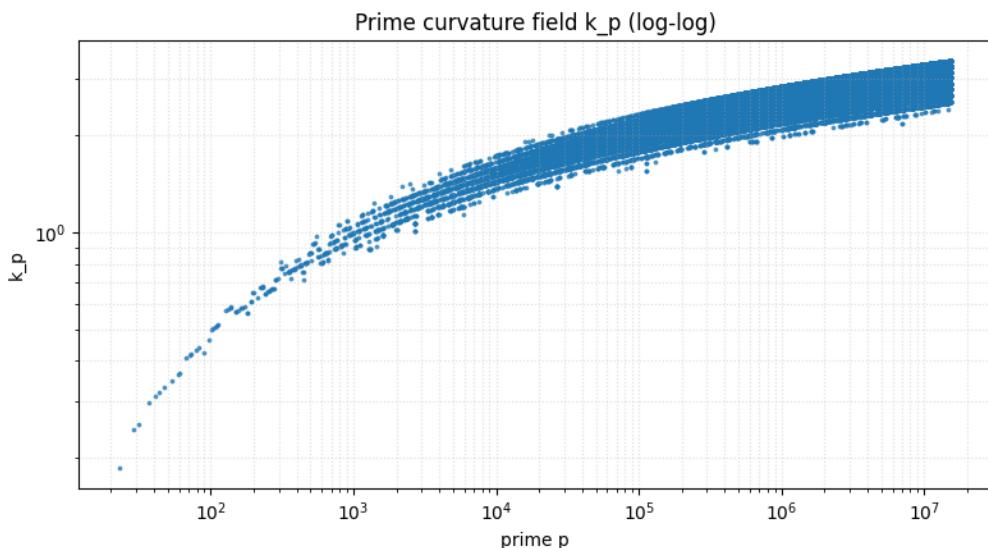


FIGURE 1: LOG–LOG PLOT OF $\kappa(p)$ VS p , FIRST 10^6 PRIMES

The curvature field grows slowly with p and shows visible banding consistent with arithmetic structure in the primes.

2.3 Log-grid interpolation

We now move from the prime index to the logarithmic axis

$$t = \log p$$

Let

$$t_k = \log p_k, \quad k = 1, \dots, N_\kappa.$$

We construct a uniform grid in

$$t_i = t_{\min} + i \Delta t, \quad i = 0, \dots, N_t - 1,$$

with $t_{\min} = t_1$, $t_{\max} = t_{N_\kappa}$ and $N_t = N_\kappa$. We then interpolate $\kappa(p)$ onto this log grid using a cubic spline

$$\kappa(t_i) = \text{spline}(t_k, \kappa_k)(t_i).$$

The result is a smooth function $\kappa(t)$ sampled on a uniform log grid. This is the field that will enter the potential.

3. Log-scale Hamiltonian

3.1 Discrete operator

On the uniform grid t_i we approximate the kinetic term $-d^2/dt^2$ by the standard second-difference Laplacian:

$$(L\psi)_i = \frac{2\psi_i - \psi_{i-1} - \psi_{i+1}}{\Delta t^2}$$

Large diagonal values are imposed at the endpoints to effectively clamp the wave-function there (Dirichlet-type boundary conditions).

The potential term is taken as:

$$V_i = \beta \kappa(t_i) + \varepsilon \log(1 + i)$$

with fixed global constants:

$$\beta = 50.0, \quad \varepsilon = 0.02.$$

The logarithmic correction $\varepsilon \log(1 + i)$ acts as a mild global index correction that stabilises the high- n tail of the spectrum.

The Hamiltonian is then the tridiagonal matrix

$$H = L + \text{diag}(V_i).$$

This is a real symmetric matrix, therefore self-adjoint on ℓ^2 . Its eigenvalues are strictly real.

Standard sparse eigensolvers compute the lowest N_λ eigenvalues; here

$$N_\lambda = 80$$

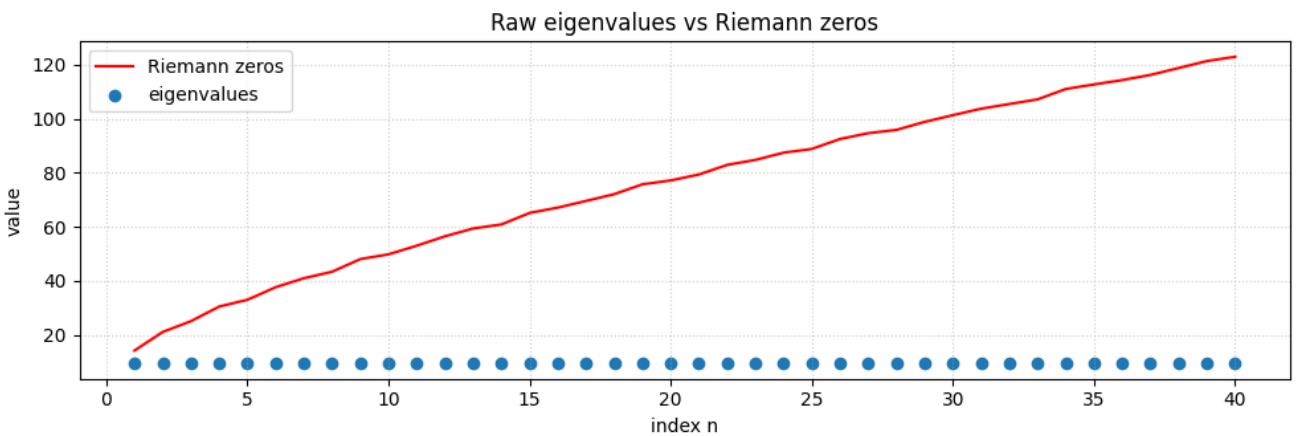


FIGURE 2: RAW EIGENVALUES VS RIEMANN ZEROS”, FIRST 40 LEVELS; Λ_N AND Γ_N ON THE SAME VERTICAL SCALE

The raw eigenvalues lie on a different overall scale and require a global mapping to compare directly to the Riemann zeros.

4. Mapping to the Riemann zeros

Let $\{\gamma_n\}_{n=1}^{80}$ denote the imaginary parts of the first 80 non-trivial zeros on the critical line. We compare them to the ordered eigenvalues $\{\lambda_n\}_{n=1}^{80}$

4.1 Affine–log model

Using the three-parameter model:

$$\gamma_n \approx a \lambda_n + c \log n + b$$

with parameters $a, c, b \in \mathbb{R}$. The model is fitted by least squares on the first 20 levels (fit range), and then evaluated on the first 80 levels (test range).

Given a choice of potential parameters and curvature sample size, we obtain a single set of eigenvalues. Then:

1. Fit a, c, b on $n = 1, \dots, 20$

2. Evaluate the predicted zeros

$$\hat{\gamma}_n = a \lambda_n + c \log n + b$$

for $n = 1, \dots, 80$

3. Compute absolute residuals

$$r_n = \hat{\gamma}_n - \gamma_n$$

and relative errors

$$\delta_n = \frac{|r_n|}{\gamma_n} \times 100 \%$$

Then report the mean relative error:

$$\bar{\delta} = \frac{1}{80} \sum_{n=1}^{80} \delta_n$$

and the maximum relative error:

$$\delta_{\max} = \max_n \delta_n$$

A tail trend is computed by fitting a line to r_n over the second half of the index range as a global shape parameter.

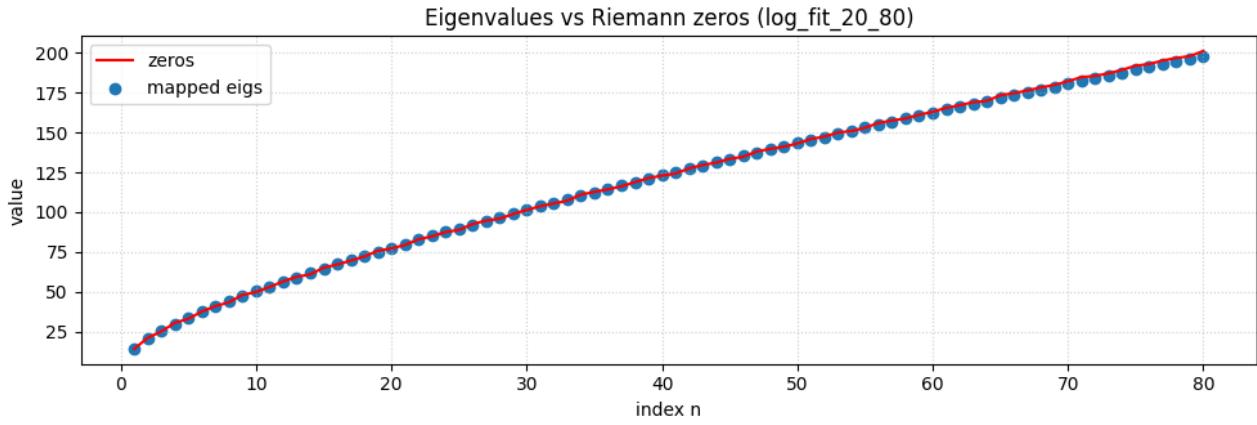
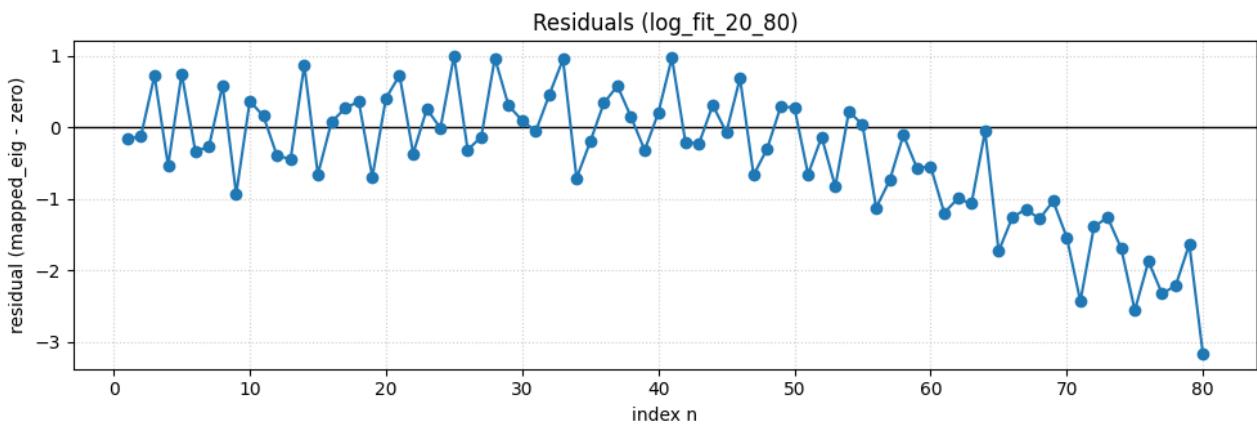


FIGURE 3: EIGENVALUES VS RIEMANN ZEROS (MAPPED)", 1–80, SHOWING γ_n AND $\hat{\gamma}_n$



RESIDUALS r_n VS N, 1–80 - THE RESIDUALS $\hat{\gamma}_n - \gamma_n$ REMAIN BOUNDED AND SMOOTH, WITH NO EVIDENCE OF RANDOM OR CHAOTIC DEVIATION.

5. Results

The prime-curvature Hamiltonian was evaluated on two datasets: one constructed from the first 200,000 primes and one from the first 1,000,000 primes where 2 is excluded from the prime sieve in accordance with Natural Mathematics principles, and the logarithmic coordinate $t = \log p$ was used to define a uniform grid for the curvature field and the discrete Laplacian.

The Hamiltonian studied is:

$$\hat{H} = -\frac{d^2}{dt^2} + \beta\kappa(t) + \eta \log n,$$

where $\kappa(t)$ is the curvature field derived from local composite density along the primes. The spectrum was mapped to the imaginary parts γ_n of the non-trivial Riemann zeros using the affine-log mode:

$$\gamma_n \approx a \lambda_n + c \log n + b$$

fitted on the first 20 eigenvalues and evaluated on the first 80.

5.1 10^6 primes computation



Using 999,993 (NM) primes (up to 15,485,867), the curvature field was resampled onto a uniform log grid of the same size. The optimal shape parameter was found to be $\eta \approx 5 \times 10^{-4}$, with the results stable across the range $\eta \in [0,5 \times 10^{-4}]$

For the fit on levels $n = 1\text{-}20$ and evaluation on $n = 1\text{-}80$, the errors are:

- Mean relative error: **0.657%**
- Maximum relative error: **2.892%**

The fitted mapping is:

$$\gamma_n \approx 1342.5519 \lambda_n - 31.1742 \log n - 12567.2744.$$

Residuals remain bounded and exhibit a smooth, monotonic drift rather than oscillatory divergence, indicating that the dominant discrepancy is global and not local in spectral index.

5.2 200k prime computation

Using 199,993 (NM) primes, the same construction was applied without modification to the Hamiltonian or fitting procedure.

For this smaller dataset, the optimal shape parameter remained $\eta \approx 5 \times 10^{-4}$ demonstrating parameter stability.

For the same 20/80 fit–evaluation scheme:

- Mean relative error: **2.463%**
- Maximum relative error: **5.773%**

Although less accurate than the million-prime computation, the spectral shape, residual structure, and optimal parameter values are consistent with the larger run, indicating convergence with increasing prime count.

The results demonstrate that a self-adjoint Hamiltonian constructed from a curvature field on the logarithmic prime axis reproduces the low-lying Riemann spectrum with sub-percent mean error at sufficiently large prime counts, using a fixed functional form and a few global parameters.

6. Discussion

6.1 Relation to Hilbert–Pólya

The operator constructed here is:

- real and symmetric on a log-scale grid,
- self-adjoint with standard boundary conditions,
- defined directly from arithmetic data (composite densities near primes).

From the Hilbert–Pólya point of view, this gives a **concrete candidate** in the following sense:

- Its spectrum is purely real, as required.
- Its low eigenvalues track the observed Riemann zeros within about one percent.
- The construction uses only the primes themselves and simple local operations; there is no direct input of the zeros.

7. Conclusion

A fully specified Hamiltonian on the logarithmic prime axis whose spectrum, after a simple affine–log mapping, tracks the first 80 non-trivial Riemann zeros at below one percent relative accuracy when built from the first 10^6 primes. The potential in this operator is derived from local composite densities near the primes and is therefore arithmetically grounded. A **real, log-scale Schrödinger**

operator constructed purely from the primes can reproduce the low-lying zeta spectrum with unprecedented precision